

```

# Step 1: Import necessary libraries
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

# Step 2: Loading the dataset
data = pd.read_csv('/content/Telco-Customer-Churn-dataset.csv')

# Step 3: Data Preprocessing

# Check for missing values
print(data.isnull().sum())

➡ customerID      0
gender            0
SeniorCitizen    0
Partner          0
Dependents       0
tenure           0
PhoneService     0
MultipleLines    0
InternetService  0
OnlineSecurity   0
OnlineBackup     0
DeviceProtection 0
TechSupport      0
StreamingTV      0
StreamingMovies  0
Contract         0
PaperlessBilling 0
PaymentMethod    0
MonthlyCharges   0
TotalCharges     0
Churn            0
dtype: int64

# Convert 'TotalCharges' to numeric and handle any spaces or non-numeric values
data['TotalCharges'] = pd.to_numeric(data['TotalCharges'], errors='coerce')

# Now that 'TotalCharges' is numeric, you may fill any remaining invalid values with the median or mean
data['TotalCharges'].fillna(data['TotalCharges'].median(), inplace=True)

# Drop the 'customerID' column as it is not useful for prediction
data.drop(['customerID'], axis=1, inplace=True)

# Encode categorical features using Label Encoding or One-Hot Encoding
label_encoder = LabelEncoder()

# List of columns to label encode (Binary and ordinal categories)
label_encode_cols = ['gender', 'Partner', 'Dependents', 'PhoneService', 'PaperlessBilling', 'Churn']

for col in label_encode_cols:
    data[col] = label_encoder.fit_transform(data[col])

# Apply One-Hot Encoding for remaining categorical columns
data = pd.get_dummies(data, columns=['MultipleLines', 'InternetService', 'OnlineSecurity', 'OnlineBackup',
                                     'DeviceProtection', 'TechSupport', 'StreamingTV', 'StreamingMovies',
                                     'Contract', 'PaymentMethod'], drop_first=True)

# Step 4: Splitting the data into train and test sets
# Define X (features) and y (target)
X = data.drop('Churn', axis=1) # 'Churn' is the target column
y = data['Churn']


# Split into train and test sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Step 5: Train models and compare

# Logistic Regression

```

```
log_model = LogisticRegression(max_iter=1000)
log_model.fit(X_train, y_train)
```

 /usr/local/lib/python3.10/dist-packages/sklearn/linear\_model/\_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
    LogisticRegression
```

```
LogisticRegression(max_iter=1000))
```

```
# Predictions and evaluation for Logistic Regression
```

```
y_pred_log = log_model.predict(X_test)
```

```
log_accuracy = accuracy_score(y_test, y_pred_log)
```

```
print("Logistic Regression Accuracy: ", log_accuracy)
```

```
print("Confusion Matrix: \n", confusion_matrix(y_test, y_pred_log))
```

```
print("Classification Report: \n", classification_report(y_test, y_pred_log))
```

 Logistic Regression Accuracy: 0.8211497515968772

Confusion Matrix:

```
[[933 103]
```

```
[149 224]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.86	0.90	0.88	1036
1	0.69	0.60	0.64	373
accuracy			0.82	1409
macro avg	0.77	0.75	0.76	1409
weighted avg	0.82	0.82	0.82	1409

```
# Decision Tree Classifier
```

```
dt_model = DecisionTreeClassifier(max_depth=5) # You can tune parameters
```

```
dt_model.fit(X_train, y_train)
```

 DecisionTreeClassifier

```
DecisionTreeClassifier(max_depth=5)
```

```
# Predictions and evaluation for Decision Tree
```


```
y_pred_dt = dt_model.predict(X_test)
```

```
dt_accuracy = accuracy_score(y_test, y_pred_dt)
```

```
print("Decision Tree Accuracy: ", dt_accuracy)
```

```
print("Confusion Matrix: \n", confusion_matrix(y_test, y_pred_dt))
```

```
print("Classification Report: \n", classification_report(y_test, y_pred_dt))
```

 Decision Tree Accuracy: 0.8062455642299503

Confusion Matrix:

```
[[964 72]
```

```
[201 172]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.83	0.93	0.88	1036
1	0.70	0.46	0.56	373
accuracy			0.81	1409
macro avg	0.77	0.70	0.72	1409
weighted avg	0.80	0.81	0.79	1409

```
# Step 6: Compare model performances
```

```
print(f"Logistic Regression Accuracy: {log_accuracy}")
```

```
print(f"Decision Tree Accuracy: {dt_accuracy}")
```

```
--INSERT--
```

 Logistic Regression Accuracy: 0.8211497515968772

Decision Tree Accuracy: 0.8062455642299503